

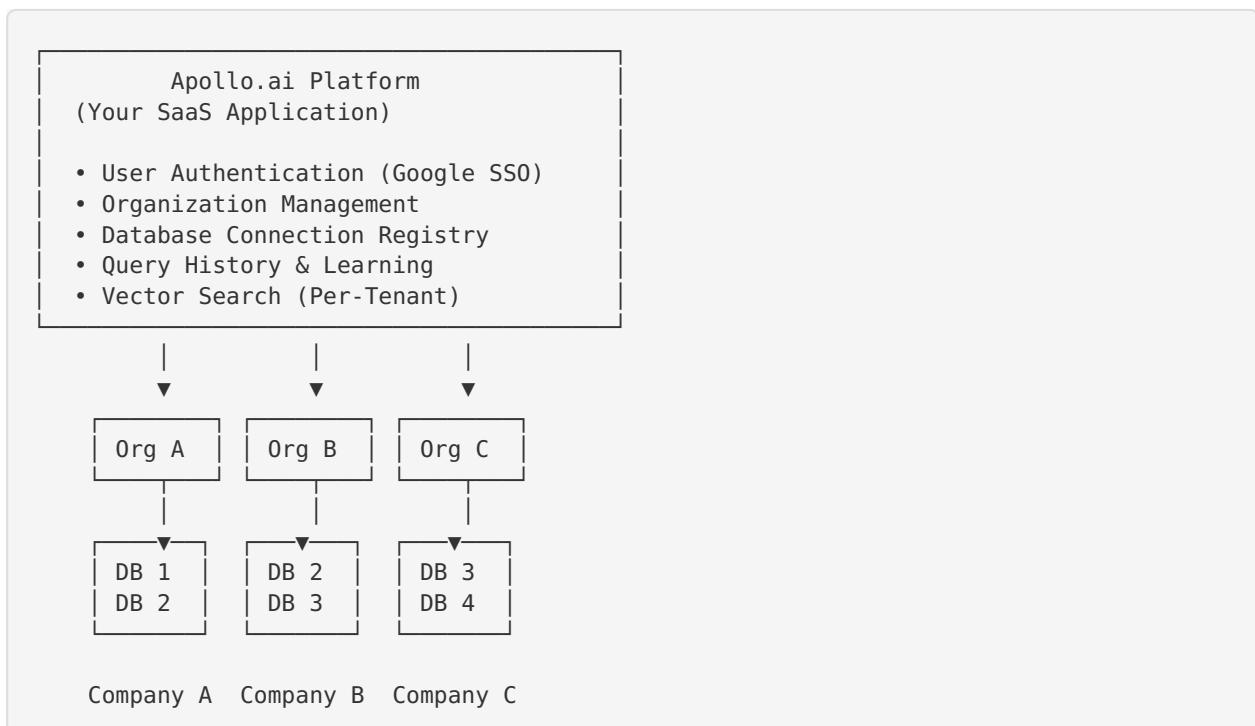
Apollo.ai Multi-Tenant SaaS Guide

🎯 Overview

Apollo.ai has been transformed into a **multi-tenant SaaS platform** that allows any business to connect their databases and start querying with natural language. This guide explains the architecture and how customers can get started.

🏗️ Architecture

How Multi-Tenancy Works



Key Components

1. **Organizations** - Each customer creates an organization
2. **Database Connections** - Organizations connect their databases
3. **Automatic Schema Discovery** - System learns each database structure
4. **Isolated Learning** - Query patterns are scoped per organization
5. **Secure Credentials** - All credentials encrypted at rest



Customer Onboarding Flow

Step 1: Sign Up

User signs up with Google SSO → Automatic organization created

Step 2: Add Database Connection

```
// Example: Adding a PostgreSQL database
{
  "name": "Production DB",
  "type": "postgresql",
  "host": "prod-db.company.com",
  "port": 5432,
  "database": "sales_db",
  "username": "readonly_user",
  "password": "encrypted_password",
  "ssl": true
}
```

Step 3: Automatic Schema Discovery

Apollo.ai automatically:

- Discovers all tables **and** columns
- Identifies relationships (foreign keys)
- Detects data types
- Caches schema **for** fast queries
- Syncs periodically **for** schema changes

Step 4: Start Querying

Natural Language → SQL → Results
 "Show me top customers by revenue" → SELECT... → Data + Visualizations

Supported Database Types

Database	Version	Status
PostgreSQL	9.6+	 Supported
MySQL	5.7+	 Supported
MariaDB	10.3+	 Supported
Oracle	11g+	 Supported
SQLite	3.x	 Supported
MongoDB	4.0+	 Coming Soon

Security Features

Per-Tenant Isolation

- Each organization's data is completely isolated
- Query history scoped by organization
- Vector embeddings tagged with `organizationId`
- No data leakage between tenants

Credential Encryption

```
// All database credentials encrypted at rest
const encryptedCredentials = encrypt({
  host, port, database, username, password, ssl
})
```

Role-Based Access Control

- **Owner:** Full control, can delete organization
- **Admin:** Manage members, databases, settings
- **Member:** Query databases, view history
- **Viewer:** Read-only access to queries and results

Schema Learning System

How It Works

1. Initial Discovery

```
sql
-- Apollo.ai runs these queries on first connection
SELECT table_name FROM information_schema.tables;
SELECT column_name, data_type FROM information_schema.columns;
SELECT constraint_name FROM information_schema.table_constraints;
```

2. Relationship Detection

```
typescript
// Automatically identifies foreign keys
{
  "foreignKeys": [
    {
      "column": "customer_id",
      "referencedTable": "customers",
      "referencedColumn": "id"
    }
  ]
}
```

3. Sample Data Collection

```
sql
```

```
-- Gets sample data for better query generation
SELECT * FROM table_name LIMIT 3;
```

4. Caching & Sync

- Schema cached in Apollo.ai database
 - Periodic refresh (daily) to catch schema changes
 - Manual refresh available via API
-

🎯 Query Intelligence (Per-Tenant)

Vector Database Integration

```
// Each query pattern is stored with organization context
{
  "organizationId": "org_abc123",
  "query": "Show me top customers",
  "sql": "SELECT * FROM customers ORDER BY total_spent DESC LIMIT 10",
  "embedding": [0.234, 0.567, ...], // Vector representation
  "confidence": 0.95
}
```

Learning Over Time

- Each organization builds its own “knowledge base”
 - Similar queries get higher confidence
 - Learns from corrections and feedback
 - No cross-contamination between organizations
-

🔧 API Reference

Organization Management

```
// Create Organization
POST /api/organizations
{
  "name": "Acme Corp"
}

// List User's Organizations
GET /api/organizations

// Get Organization Details
GET /api/organizations/{orgId}

// Add Team Member
POST /api/organizations/{orgId}/members
{
  "email": "colleague@company.com",
  "role": "MEMBER"
}
```

Database Connections

```
// Add Database Connection
POST /api/database-connections
{
  "organizationId": "org_abc123",
  "name": "Production DB",
  "type": "postgresql",
  "host": "db.company.com",
  "port": 5432,
  "database": "sales",
  "username": "readonly",
  "password": "secure_pass",
  "ssl": true
}

// Test Connection
POST /api/database-connections
{
  "action": "test",
  "organizationId": "org_abc123",
  ...connectionConfig
}

// List Connections
GET /api/database-connections?organizationId={orgId}

// Remove Connection
DELETE /api/database-connections?id={connectionId}
```

Query Execution

```
// Execute Query
POST /api/query
{
  "organizationId": "org_abc123",
  "query": "Show me sales by region",
  "database": "production_db"
}

// Response includes:
{
  "sql": "SELECT region, SUM(amount) ...",
  "results": [...],
  "summary": "Sales data grouped by region",
  "confidence": 0.92,
  "visualization": {...}
}
```

Pricing Plans (Example)

Plan	Price	Features
Free	\$0/mo	1 database, 100 queries/mo
Starter	\$49/mo	3 databases, 1,000 queries/mo
Professional	\$199/mo	10 databases, 10,000 queries/mo
Enterprise	Custom	Unlimited databases, custom SLA

Migration from Single-Tenant

If you have existing Apollo.ai installation:

```
# Run migration script
cd nextjs_space
yarn tsx scripts/migrate-to-multitenant.ts

# Results:
✓ Organizations created for all users
✓ All existing data linked to organizations
✓ System now multi-tenant ready!
```

Scaling Considerations

Database Connection Pooling

- Each organization gets dedicated connection pools
- Configurable max connections per database
- Automatic cleanup of idle connections

Query Performance

- Schema cache reduces database roundtrips
- Vector search for fast query matching
- Result caching (optional)

Storage

- Query history: ~1KB per query
- Schema cache: ~100KB per database
- Vector embeddings: ~2KB per query pattern

UI Components (To Be Built)

Onboarding Wizard

```
Step 1: Create Organization → "Acme Corp"
Step 2: Add Database → Connection form
Step 3: Test Connection →  Success
Step 4: Discover Schema → Automatic
Step 5: Start Querying → Dashboard
```

Organization Switcher

[Dropdown in header]

- └ Acme Corp (Current)
- └ Personal Workspace
- └ + New Organization

Database Management

Connected Databases

<input type="checkbox"/>	Production DB (PostgreSQL)	<input type="checkbox"/>
<input type="checkbox"/>	5 tables	<input checked="" type="checkbox"/> Last sync: 2h ago
<input type="checkbox"/>	[Refresh] [Edit] [Remove]	

Testing Multi-Tenancy

```
# Create two test organizations
curl -X POST /api/organizations \
-H "Authorization: Bearer TOKEN_USER_1" \
-d '{"name": "Company A"}'

curl -X POST /api/organizations \
-H "Authorization: Bearer TOKEN_USER_2" \
-d '{"name": "Company B"}'

# Add different databases to each
# Verify queries are isolated
# Confirm no data leakage
```

Support & Documentation

- **API Docs:** [/docs/api](#)
- **Schema Discovery:** [/docs/schema-discovery](#)
- **Security:** [/docs/security](#)
- **Compliance:** [/docs/compliance](#)



Benefits of Multi-Tenant Apollo.ai

- Easy Installation** - Just connect your database
 - No Code Required** - Natural language interface
 - Secure** - Enterprise-grade encryption
 - Scalable** - Supports multiple databases per org
 - Smart** - Learns your database over time
 - Compliant** - GDPR, SOC2 ready
 - Fast** - Cached schemas, optimized queries
-



Roadmap

- [] Database connection wizard UI
 - [] Organization switcher component
 - [] Onboarding flow
 - [] Billing integration (Stripe)
 - [] Advanced role permissions
 - [] Database connection health monitoring
 - [] Schema change notifications
 - [] Query optimization suggestions
 - [] MongoDB support
 - [] Real-time collaboration
-

Apollo.ai - Making databases accessible to everyone, everywhere.